Grammar of data wrangling
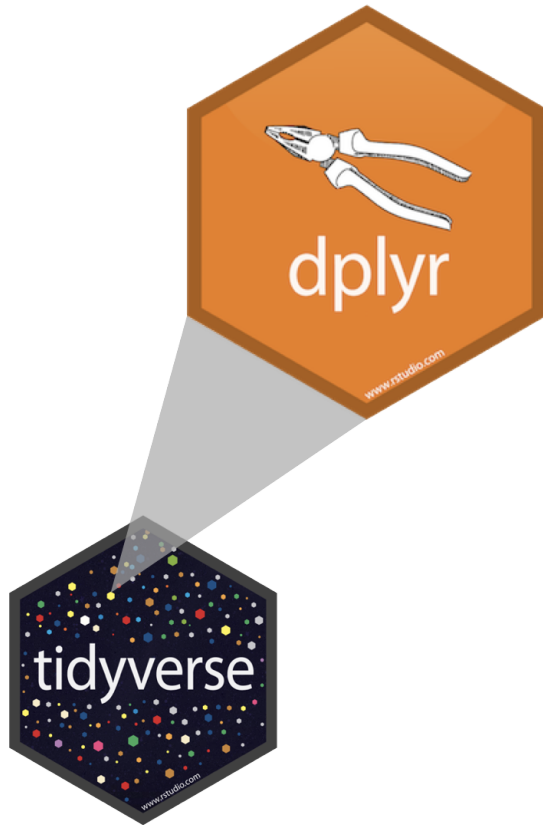
Learn how to transfrom data
Prof. Dr. Jan Kirenz

The following content is based on Mine Çetinkaya-Rundel's excellent material in "Data Science in a Box"

Grammar of data wrangling

A grammar of data wrangling...

... based on the concepts of functions as verbs that manipulate data frames

- `select`: pick columns by name
- `arrange`: reorder rows
- `slice`: pick rows using index(es)
- `filter`: pick rows matching criteria
- `distinct`: filter for unique rows
- `mutate`: add new variables
- `summarise`: reduce variables to values
- `group_by`: for grouped operations
- ... (many more)

Rules of dplyr functions

- First argument is *always* a data frame

- Subsequent arguments say what to do with that data frame

- Always return a data frame

- Don't modify in place

Data: **Hotel bookings**

- Data from two hotels: one resort and one city hotel

- Observations: Each row represents a hotel booking

- Goal for original data collection: Development of prediction models to classify a hotel booking's likelihood to be cancelled (Antonia et al., 2019)

```
hotels <- read_csv("data/hotels.csv")
```

Source: TidyTuesday

First look: Variables

```
names(hotels)
```

```
##  [1] "hotel"
##  [2] "is_canceled"
##  [3] "lead_time"
##  [4] "arrival_date_year"
##  [5] "arrival_date_month"
##  [6] "arrival_date_week_number"
##  [7] "arrival_date_day_of_month"
##  [8] "stays_in_weekend_nights"
##  [9] "stays_in_week_nights"
## [10] "adults"
## [11] "children"
## [12] "babies"
## [13] "meal"
## [14] "country"
## [15] "market_segment"
## [16] "distribution_channel"
## [17] "is_repeated_guest"
## [18] "previous_cancellations"
...
```

Second look: Overview

```
glimpse(hotels)
```

```
## Rows: 119,390
## Columns: 32
## $ hotel                        <chr> "Resort Hotel", "Resort…
## $ is_canceled                  <dbl> 0, 0, 0, 0, 0, 0, 0, 0,…
## $ lead_time                    <dbl> 342, 737, 7, 13, 14, 14…
## $ arrival_date_year            <dbl> 2015, 2015, 2015, 2015,…
## $ arrival_date_month           <chr> "July", "July", "July",…
## $ arrival_date_week_number     <dbl> 27, 27, 27, 27, 27, 27,…
## $ arrival_date_day_of_month    <dbl> 1, 1, 1, 1, 1, 1, 1, 1,…
## $ stays_in_weekend_nights      <dbl> 0, 0, 0, 0, 0, 0, 0, 0,…
## $ stays_in_week_nights         <dbl> 0, 0, 1, 1, 2, 2, 2, 2,…
## $ adults                       <dbl> 2, 2, 1, 1, 2, 2, 2, 2,…
## $ children                     <dbl> 0, 0, 0, 0, 0, 0, 0, 0,…
## $ babies                       <dbl> 0, 0, 0, 0, 0, 0, 0, 0,…
## $ meal                         <chr> "BB", "BB", "BB", "BB",…
## $ country                      <chr> "PRT", "PRT", "GBR", "G…
## $ market_segment               <chr> "Direct", "Direct", "Di…
## $ distribution_channel         <chr> "Direct", "Direct", "Di…
...
```

Select a single column

View only `lead_time` (number of days between booking and arrival date):

```
select(hotels, lead_time)
```

```
## # A tibble: 119,390 x 1
##    lead_time
##        <dbl>
## 1        342
## 2        737
## 3          7
## 4         13
## 5         14
## 6         14
## # … with 119,384 more rows
```

Select a single column

```
select(
  hotels,
  lead_time
)
```

- Start with the function (a verb): `select()`

Select a single column

```
select(
  hotels,
  lead_time
  )
```

- Start with the function (a verb): `select()`
- First argument: data frame we're working with , `hotels`

Select a single column

```
select(
  hotels,
  lead_time
)
```

- Start with the function (a verb): `select()`
- First argument: data frame we're working with , `hotels`
- Second argument: variable we want to select, `lead_time`

Select a single column

```
select(
  hotels,
  lead_time
  )
```

```
## # A tibble: 119,390 x 1
##   lead_time
##       <dbl>
## 1       342
## 2       737
## 3         7
## 4        13
## 5        14
## 6        14
## # … with 119,384 more rows
```

- Start with the function (a verb): `select()`
- First argument: data frame we're working with , `hotels`
- Second argument: variable we want to select, `lead_time`
- Result: data frame with 119390 rows and 1 column

dplyr functions always expect a data frame and always yield a data frame.

```
select(hotels, lead_time)
```

```
## # A tibble: 119,390 x 1
##    lead_time
##        <dbl>
## 1        342
## 2        737
## 3          7
## 4         13
## 5         14
## 6         14
## # … with 119,384 more rows
```

Select multiple columns

View only the `hotel` type and `lead_time`:

```
select(hotels, hotel, lead_time)
```

```
## # A tibble: 119,390 x 2
##   hotel         lead_time
##   <chr>             <dbl>
## 1 Resort Hotel        342
## 2 Resort Hotel        737
## 3 Resort Hotel          7
## 4 Resort Hotel         13
## 5 Resort Hotel         14
## 6 Resort Hotel         14
## # … with 119,384 more rows
```

What if we wanted to select these columns, and then arrange the data in descending order of lead time?

Data wrangling, step-by-step

Select:

```
hotels %>%
  select(hotel, lead_time)
```

```
## # A tibble: 119,390 x 2
##   hotel         lead_time
##   <chr>             <dbl>
## 1 Resort Hotel        342
## 2 Resort Hotel        737
## 3 Resort Hotel          7
## 4 Resort Hotel         13
## 5 Resort Hotel         14
## 6 Resort Hotel         14
## # … with 119,384 more rows
```

Select, then arrange:

```
hotels %>%
  select(hotel, lead_time) %>%
  arrange(desc(lead_time))
```

```
## # A tibble: 119,390 x 2
##   hotel         lead_time
##   <chr>             <dbl>
## 1 Resort Hotel        737
## 2 Resort Hotel        709
## 3 City Hotel          629
## 4 City Hotel          629
## 5 City Hotel          629
## 6 City Hotel          629
## # … with 119,384 more rows
```

Pipes

What is a pipe?

In programming, a pipe is a technique for passing information from one process to another.

- Start with the data frame `hotels`, and pass it to the `select()` function,

```
hotels %>%
  select(hotel, lead_time) %>%
  arrange(desc(lead_time))
```

```
## # A tibble: 119,390 x 2
##   hotel        lead_time
##   <chr>            <dbl>
## 1 Resort Hotel       737
## 2 Resort Hotel       709
## 3 City Hotel         629
## 4 City Hotel         629
## 5 City Hotel         629
## 6 City Hotel         629
## # … with 119,384 more rows
```

What is a pipe?

In programming, a pipe is a technique for passing information from one process to another.

- Start with the data frame `hotels`, and pass it to the `select()` function,
- then we select the variables `hotel` and `lead_time`,

```
hotels %>%
  select(hotel, lead_time) %>%
  arrange(desc(lead_time))
```

```
## # A tibble: 119,390 x 2
##   hotel        lead_time
##   <chr>             <dbl>
## 1 Resort Hotel        737
## 2 Resort Hotel        709
## 3 City Hotel          629
## 4 City Hotel          629
## 5 City Hotel          629
## 6 City Hotel          629
## # … with 119,384 more rows
```

What is a pipe?

In programming, a pipe is a technique for passing information from one process to another.

- Start with the data frame `hotels`, and pass it to the `select()` function,
- then we select the variables `hotel` and `lead_time`,
- and then we arrange the data frame by `lead_time` in descending order.

```
hotels %>%
  select(hotel, lead_time) %>%
  arrange(desc(lead_time))
```

```
## # A tibble: 119,390 x 2
##   hotel          lead_time
##   <chr>               <dbl>
## 1 Resort Hotel          737
## 2 Resort Hotel          709
## 3 City Hotel            629
## 4 City Hotel            629
## 5 City Hotel            629
## 6 City Hotel            629
## # … with 119,384 more rows
```

Aside

The pipe operator is implemented in the package **magrittr**, though we don't need to load this package explicitly since **tidyverse** does this for us.

Any guesses as to why the package is called magrittr?

Ceci n'est pas une pipe.

%>%
magrittr

Ceci n'est pas un pipe.

How does a pipe work?

- You can think about the following sequence of actions - find keys, unlock car, start car, drive to work, park.
- Expressed as a set of nested functions in R pseudocode this would look like:

```r
park(drive(start_car(find("keys")), to = "work"))
```

- Writing it out using pipes give it a more natural (and easier to read) structure:

```r
find("keys") %>%
  start_car() %>%
  drive(to = "work") %>%
  park()
```

A note on piping and layering

- %>% used mainly in **dplyr** pipelines, *we pipe the output of the previous line of code as the first input of the next line of code*
- + used in **ggplot2** plots is used for "layering", *we create the plot in layers, separated by* +

dplyr

❌

```
hotels +
  select(hotel, lead_time)
```

Error in select(hotel, lead_time): Objekt 'hotel' nicht gefunden

✅

```
hotels %>%
  select(hotel, lead_time)
```

```
## # A tibble: 119,390 x 2
##   hotel         lead_time
##   <chr>             <dbl>
## 1 Resort Hotel        342
## 2 Resort Hotel        737
## 3 Resort Hotel          7
...
```
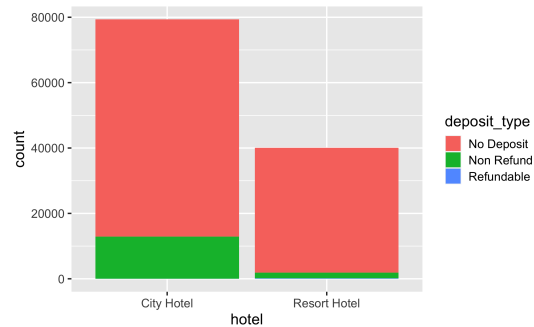
ggplot2

❌

```
ggplot(hotels, aes(x = hotel, fill = deposit_type)) %>%
  geom_bar()
```

```
## Error: `mapping` must be created by `aes()`
## Did you use %>% instead of +?
```

✅

```
ggplot(hotels, aes(x = hotel, fill = deposit_type)) +
  geom_bar()
```

Code styling

Many of the styling principles are consistent across %>% and +:

- always a space before
- always a line break after (for pipelines with more than 2 lines)

❌

```
ggplot(hotels,aes(x=hotel,y=deposit_type))+geom_bar()
```

✅

```
ggplot(hotels, aes(x = hotel, y = deposit_type)) +
  geom_bar()
```