

Model building

Cross validation

Prof. Dr. Jan Kirenz

The following content is based on Mine Çetinkaya-Rundel's excellent book Data Science in a Box

Data and exploration

the office



Data

```
office_ratings <- read_csv("data/office_ratings.csv")
office_ratings
```

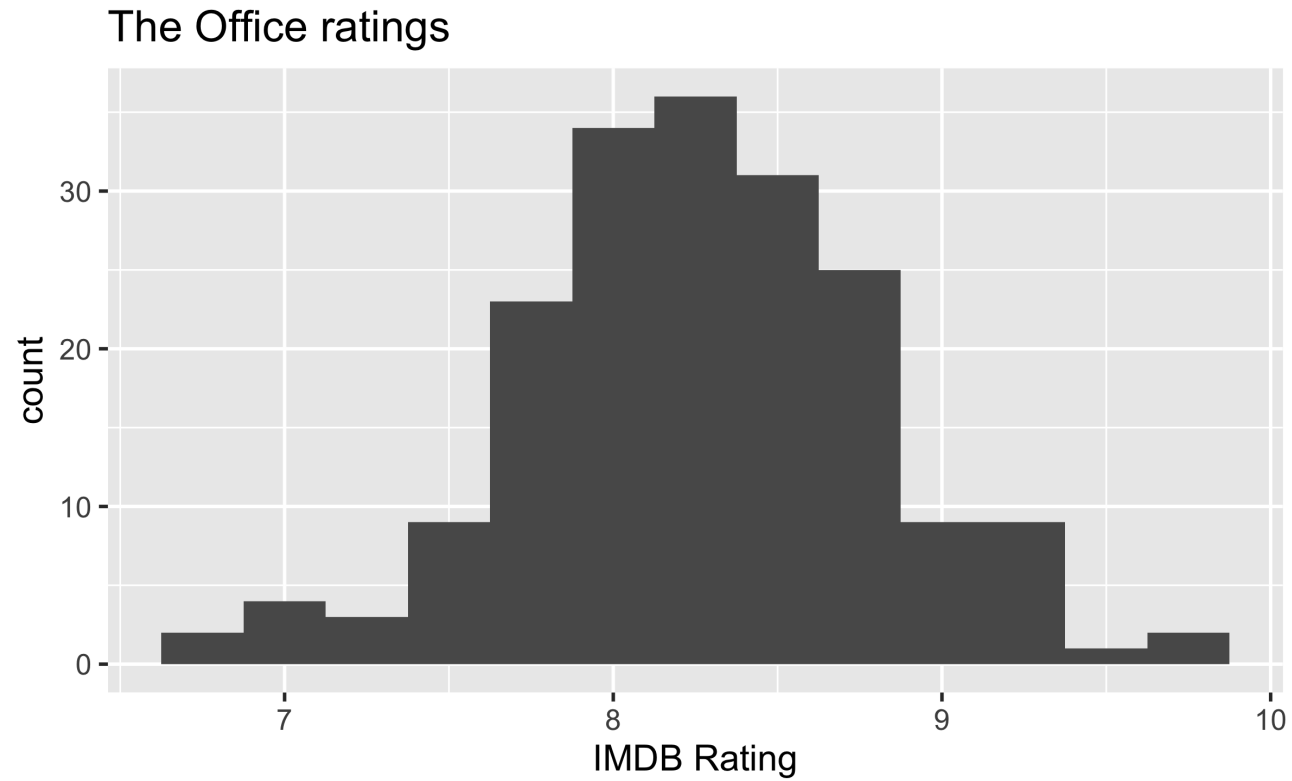
```
## # A tibble: 188 x 6
##   season episode title          imdb_rating total_votes air_date
##   <dbl>   <dbl> <chr>          <dbl>         <dbl> <date>
## 1     1     1     1 Pilot             7.6           3706 2005-03-24
## 2     1     2 Diversity Day       8.3           3566 2005-03-29
## 3     1     3 Health Care        7.9           2983 2005-04-05
## 4     1     4 The Alliance       8.1           2886 2005-04-12
## 5     1     5 Basketball          8.4           3179 2005-04-19
## 6     1     6 Hot Girl            7.8           2852 2005-04-26
## # ... with 182 more rows
```

Source: The data come from [data.world](#), by way of [TidyTuesday](#).

IMDB ratings

Code

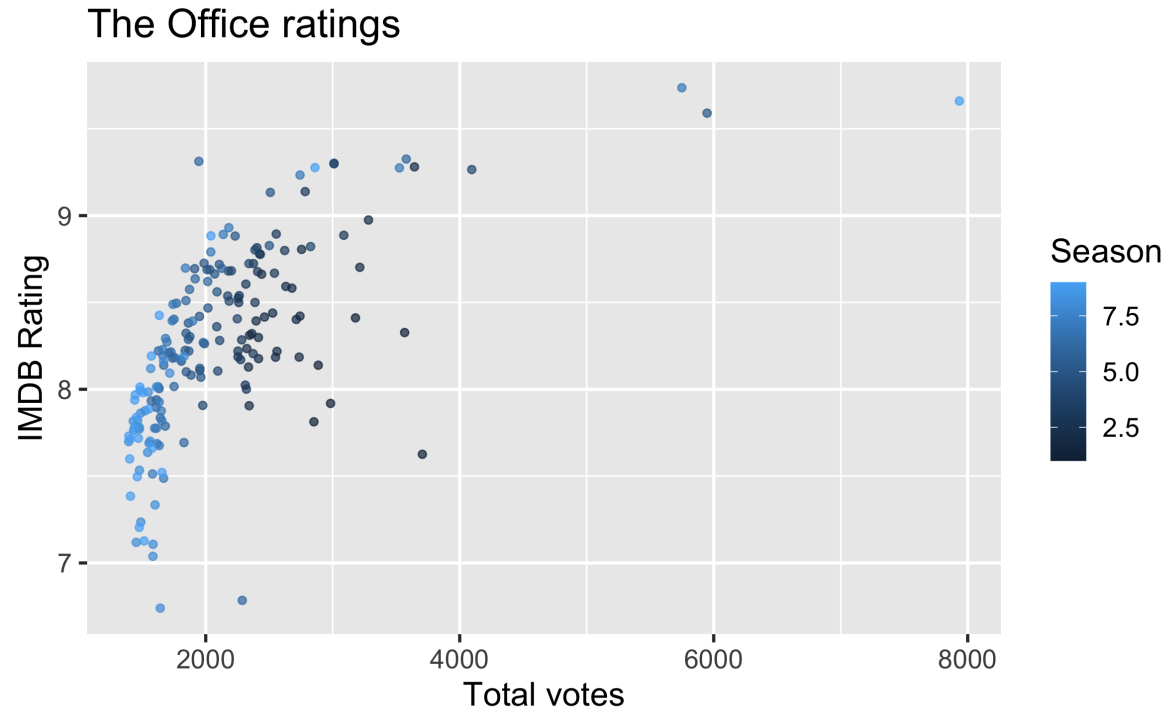
Plot



IMDB ratings vs. number of votes

Code

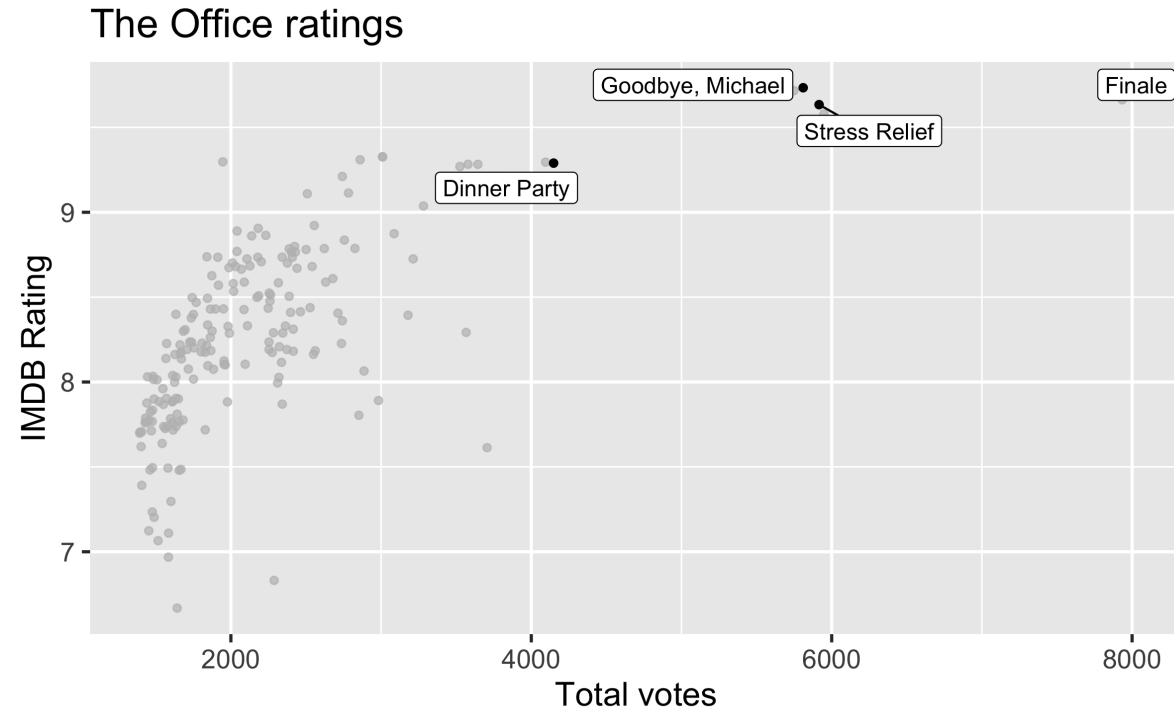
Plot



Outliers

Code

Plot




If you like the [Dinner Party](#) episode, I highly recommend this ["oral history"](#) of the episode published on Rolling Stone magazine.

IMDB ratings vs. seasons

Code

Plot

The Office ratings



Modeling

Train / test

- Create an initial split

```
set.seed(1122)  
office_split <- initial_split(office_ratings) # prop = 3/4 by default
```

- Save training data

```
office_train <- training(office_split)  
dim(office_train)
```

```
## [1] 141  6
```

- Save testing data

```
office_test <- testing(office_split)  
dim(office_test)
```

```
## [1] 47  6
```

Specify model

```
office_mod <- linear_reg() %>%  
  set_engine("lm")
```

```
office_mod
```

```
## Linear Regression Model Specification (regression)  
##  
## Computational engine: lm
```

Build recipe

Code

Output

office_rec

```
## Data Recipe
##
## Inputs:
##
##      role #variables
##      ID      1
## outcome      1
## predictor     4
##
## Operations:
##
## Date features from air_date
## Delete terms air_date
## Dummy variables from contains("month")
## Zero variance filter on all_predictors()
```

Build workflow

Code

Output

office_wflow

```
## == Workflow ==  
## Preprocessor: Recipe  
## Model: linear_reg()  
##  
## — Preprocessor —  
## 4 Recipe Steps  
##  
## ● step_date()  
## ● step_rm()  
## ● step_dummy()  
## ● step_zv()  
##  
## — Model —  
## Linear Regression Model Specification (regression)  
##  
## Computational engine: lm
```

Fit model

Code

Output

```
tidy(officer_fit) %>%  
  print(n = 12)
```

```
## # A tibble: 12 x 5  
##   term                estimate std.error statistic  p.value  
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept)        6.63      0.271     24.5 2.60e-50  
## 2 season             -0.0297   0.0180    -1.65 1.02e- 1  
## 3 episode            0.0453   0.00956    4.74 5.68e- 6  
## 4 total_votes        0.000510 0.0000692    7.38 1.75e-11  
## 5 air_date_month_Feb  0.00327  0.142      0.0229 9.82e- 1  
## 6 air_date_month_Mär -0.0585   0.146    -0.402 6.89e- 1  
## 7 air_date_month_Apr -0.150    0.141    -1.06 2.90e- 1  
## 8 air_date_month_Mai  0.0231   0.178     0.129 8.97e- 1  
## 9 air_date_month_Sep  0.646    0.180     3.58 4.80e- 4  
## 10 air_date_month_Okt 0.572    0.154     3.71 3.11e- 4  
## 11 air_date_month_Nov 0.349    0.140     2.49 1.41e- 2  
## 12 air_date_month_Dez 0.516    0.158     3.27 1.39e- 3
```

Evaluate model

Make predictions for training data

```
office_train_pred <- predict(office_fit, office_train) %>%  
  bind_cols(office_train %>% select(imdb_rating, title))
```

```
office_train_pred
```

```
## # A tibble: 141 x 3  
##   .pred imdb_rating title  
##   <dbl>      <dbl> <chr>  
## 1  8.48        7.6 Pilot  
## 2  8.45        8.3 Diversity Day  
## 3  8.10        8.1 The Alliance  
## 4  8.30        8.4 Basketball  
## 5  8.18        7.8 Hot Girl  
## 6  8.90        8.7 The Dundies  
## # ... with 135 more rows
```


R-squared

Percentage of variability in the IMDB ratings explained by the model

```
rsq(office_train_pred, truth = imdb_rating, estimate = .pred)
```

```
## # A tibble: 1 x 3  
##   .metric .estimator .estimate  
##   <chr>    <chr>         <dbl>  
## 1 rsq      standard        0.533
```

Are models with high or low R^2 more preferable?

RMSE

An alternative model performance statistic: **root mean square error**

$$RMSE = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}}$$

```
rmse(office_train_pred, truth = imdb_rating, estimate = .pred)
```

```
## # A tibble: 1 x 3  
##   .metric .estimator .estimate  
##   <chr>   <chr>      <dbl>  
## 1 rmse    standard    0.350
```

Are models with high or low RMSE are more preferable?

Interpreting RMSE

Is this RMSE considered low or high?

```
rmse(office_train_pred, truth = imdb_rating, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      0.350
```

```
office_train %>%
  summarise(min = min(imdb_rating), max = max(imdb_rating))
```

```
## # A tibble: 1 x 2
##   min    max
##   <dbl> <dbl>
## 1  6.8   9.7
```

but, really, who cares about predictions on **training** data?

Make predictions for testing data

```
office_test_pred <- predict(office_fit, office_test) %>%  
  bind_cols(office_test %>% select(imdb_rating, title))
```

```
office_test_pred
```

```
## # A tibble: 47 x 3  
##   .pred imdb_rating title  
##   <dbl>      <dbl> <chr>  
## 1  8.11         7.9 Health Care  
## 2  8.68         8.2 Halloween  
## 3  8.36         8.3 The Secret  
## 4  8.56         8.1 Michael's Birthday  
## 5  8.47         8   Grief Counseling  
## 6  8.49         8.2 Initiation  
## # ... with 41 more rows
```

Evaluate performance on testing data

- RMSE of model fit to testing data

```
rmse(office_test_pred, truth = imdb_rating, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>    <chr>         <dbl>
## 1 rmse      standard       0.482
```

- R^2 of model fit to testing data

```
rsq(office_test_pred, truth = imdb_rating, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>    <chr>         <dbl>
## 1 rsq      standard       0.496
```

Training vs. testing

metric	train	test	comparison
RMSE	0.350	0.482	RMSE lower for training
R-squared	0.533	0.496	R-squared higher for training

Evaluating performance on training data

- The training set does not have the capacity to be a good arbiter of performance.
- It is not an independent piece of information; predicting the training set can only reflect what the model already knows.
- Suppose you give a class a test, then give them the answers, then provide the same test. The student scores on the second test do not accurately reflect what they know about the subject; these scores would probably be higher than their results on the first test.

Cross validation

Cross validation

More specifically, **v-fold cross validation**:

- Shuffle your data v partitions
- Use 1 partition for validation, and the remaining v-1 partitions for training
- Repeat v times

You might also heard of this referred to as k-fold cross validation.

Cross validation

Split data into folds

```
set.seed(345)

folds <- vfold_cv(office_train, v = 5)
folds
```

```
## # 5-fold cross-validation
## # A tibble: 5 x 2
##   splits          id
##   <list>        <chr>
## 1 <split [112/29]> Fold1
## 2 <split [113/28]> Fold2
## 3 <split [113/28]> Fold3
## 4 <split [113/28]> Fold4
## 5 <split [113/28]> Fold5
```

	training					testing
fold 1	validate	train	train	train	train	
fold 2	train	validate	train	train	train	
fold 3	train	train	validate	train	train	
fold 4	train	train	train	validate	train	
fold 5	train	train	train	train	validate	

Fit resamples

```
set.seed(456)

office_fit_rs <- office_wflow %>%
  fit_resamples(folds)

office_fit_rs
```

```
## # Resampling results
## # 5-fold cross-validation
## # A tibble: 5 x 4
##   splits          id      .metrics      .no
##   <list>         <chr> <list>      <list>
## 1 <split [112/29]> Fold1 <tibble [2 x 4]> <tibble [0 x 1]>
## 2 <split [113/28]> Fold2 <tibble [2 x 4]> <tibble [0 x 1]>
## 3 <split [113/28]> Fold3 <tibble [2 x 4]> <tibble [0 x 1]>
## 4 <split [113/28]> Fold4 <tibble [2 x 4]> <tibble [0 x 1]>
## 5 <split [113/28]> Fold5 <tibble [2 x 4]> <tibble [0 x 1]>
```



Collect CV metrics

```
collect_metrics(office_fit_rs)
```

```
## # A tibble: 2 x 6
##   .metric .estimator mean      n std_err .config
##   <chr>   <chr>      <dbl> <int>   <dbl> <chr>
## 1 rmse    standard    0.378     5  0.0260 Preprocessor1_Model1
## 2 rsq     standard    0.442     5  0.0816 Preprocessor1_Model1
```

Deeper look into CV metrics

Raw

Tidy

Fold	RMSE	R-squared
Fold1	0.399	0.449
Fold2	0.306	0.749
Fold3	0.438	0.390
Fold4	0.419	0.280
Fold5	0.327	0.344

How does RMSE compare to y?

- Cross validation RMSE stats

```
## # A tibble: 1 x 6
##   min    max  mean   med    sd   IQR
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.306 0.438 0.378 0.399 0.0581 0.0918
```

- Training data IMDB score stats

```
## # A tibble: 1 x 6
##   min    max  mean   med    sd   IQR
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  6.8   9.7  8.26  8.3  0.514 0.700
```


What's next?